1. What GovernHQ Is

GovernHQ is a governance layer for AI agents. It sits between an agent's reasoning and its action. Before any AI agent acts — queries a database, calls an API, makes a decision — the agent's reasoning passes through GovernHQ. GovernHQ evaluates the thought: what does the agent want to do, why does it want to do it, and does that reasoning comply with policy? If yes, the reasoning becomes an action. If no, the reasoning stays a thought — the action never happens.

Most AI security tools intercept actions or requests at the network layer — after the agent has already decided to act. GovernHQ intercepts the reasoning — the step before the action exists. The agent keeps running. Only the bad action is stopped.

One sentence: GovernHQ intercepts AI agent reasoning and stops bad thoughts from becoming bad actions.

2. Why GovernHQ Exists

The Problem

Enterprises are deploying fleets of AI agents that operate autonomously — processing insurance claims, managing customer data, making billing decisions, calling external APIs. These agents are powerful, but they have no governance layer.

Current solutions watch what agents did after the fact:

- Agent leaks customer data → security tool detects it 4 minutes later → incident response begins
- Agent overrides a billing limit → monitoring tool flags it → damage already done
- Agent calls an unauthorized API → log shows it happened → data already sent

These are security cameras. You review the footage after the crime.

The Shift

GovernHQ moves AI agent security from detection to prevention. From security cameras to security guards.

Other tools: Agent reasons → Agent acts → Tool detects → Responds
GovernHQ: Agent reasons → GovernHQ evaluates reasoning → Allow or Block → Then agent acts (or doesn't)

- Agent reasons about leaking customer data → GovernHQ evaluates the reasoning → action blocked → data never leaves → agent keeps running
- Agent reasons about overriding a billing limit → GovernHQ checks policy → action blocked → override never happens → agent continues other work
- Agent reasons about calling unauthorized API → GovernHQ checks scope → action blocked → call never made → agent moves to next task

The reasoning never becomes an action. The damage never happens. The agent stays operational for everything else it's allowed to do.

Why Now

The AI agent market is exploding. Every major enterprise is building or deploying autonomous agents. But governance hasn't kept up:

- No standard exists for how to govern agent reasoning
- Existing tools (SIEMs, API gateways, IAM) govern the access layer, not the reasoning layer
- Regulatory pressure is building — enterprises need auditable proof that their agents are governed
- The attack surface is new — prompt injection, reasoning loops, scope creep — traditional security tools don't detect these

GovernHQ is built specifically for this gap.

3. Who GovernHQ Is For

Primary Users

AI Platform Teams at mid-to-large enterprises who are deploying fleets of autonomous agents and need governance, compliance, and control.

Security & Compliance Teams who need auditable proof that AI agents operate within policy.

Engineering Leaders who need to ship AI agents to production but can't without a safety layer.

Industries

Any industry deploying autonomous AI agents with access to sensitive data or critical operations:

- Insurance — claims processing agents handling PII
- Financial Services — trading, billing, fraud detection agents
- Healthcare — patient data, diagnostic, scheduling agents
- Enterprise SaaS — customer support, data analysis, reporting agents
- Government — document processing, compliance agents

Buyer Profile

- Companies running 10+ autonomous AI agents
- Regulated industries with audit requirements
- Teams that have built agents but are hesitant to give them production access
- CISOs and CTOs who have been asked "how do we govern our AI agents?" and don't have an answer

4. How GovernHQ Works

The Core Concept

Every AI agent's reasoning passes through GovernHQ before it becomes an action. When an agent decides it wants to do something — query a database, call an API, make a decision — it first tells GovernHQ what it's thinking and why. GovernHQ evaluates that reasoning: is the intent safe, is it in scope, is it policy-compliant? If the reasoning is sound, the agent is allowed to act — the reasoning becomes an action. If the reasoning violates policy, the action is blocked before it ever executes — the reasoning stays a thought. The agent never touched the database. The API call was never made. The thought didn't become an action. And the agent keeps running.

This is what makes GovernHQ different from every other security tool. Other tools intercept actions or requests at the network or API layer — after the agent has already decided to act. GovernHQ intercepts the reasoning — the step before the action exists.

The Invisible Infrastructure

Two systems run behind the product. Users never interact with them directly. They see the results across all three features.

The Gate

The gate is the core evaluation engine. It evaluates individual actions by intercepting the agent's reasoning before the action happens. Every time an agent wants to act, the agent submits its reasoning to the gate through the GovernHQ SDK:

```
What do I want to do?     "Query the claims database"
Why?                      "Daily triage of pending claims"
What resource?            "claims_db"
What data?                "customer_name, claim_amount"
```

The gate evaluates this reasoning against the agent's rules:

- Scope check — Is this agent allowed to access claims_db?
- Policy check — Does this reasoning lead to a policy violation?
- PII check — Is there sensitive data involved, and is this agent cleared for it?
- Safety check — Does this reasoning look like prompt injection or manipulation?
- Rate check — Is this agent making an unusual number of requests?

The gate then decides what happens to the action:

- 🟢 Action Allowed (green) — Reasoning is sound. Action executes. Agent keeps running. Logged to Ledger.
- 🟡 Action Paused (yellow) — Reasoning is ambiguous. Action waits for human review. Agent keeps running and can do other things. Logged to Ledger.
- 🔴 Action Blocked (red) — Reasoning violates rules. Action never executes. Reasoning stays a thought. Agent keeps running and can attempt other actions. Logged to Ledger.

Critical distinction: The gate blocks actions, not agents. An agent can have one action blocked and its very next action allowed. A blocked action doesn't change the agent's status. The agent stays operational. Only the specific bad action is stopped.

Average gate evaluation time: ~10 milliseconds. The agent barely notices.

The Monitor

The monitor watches agent behavior over time by analyzing patterns across all gate decisions in the Ledger. Unlike the gate, which evaluates individual reasoning and blocks individual actions, the monitor evaluates the agent itself:

- Is an agent repeating the same reasoning dozens of times? (reasoning loop)
- Is an agent's trust score declining steadily? (behavioral drift)
- Is an agent suddenly reasoning about resources it rarely touches? (anomaly)

- Is an agent's action block rate spiking? (possible compromise)

When the monitor detects a concerning pattern, it blocks the entire agent — not just one action, but the agent itself — which surfaces in Shield for human review. A blocked agent cannot perform any actions until a human reviews and allows it back.

Critical distinction:

```
Gate:      evaluates REASONING     blocks ACTIONS      agent keeps running
Monitor:   evaluates PATTERNS      blocks AGENTS       agent stops entirely
```

The gate is precise — it evaluates one thought, blocks one bad action, and lets the agent continue thinking and acting on everything else. The monitor is decisive — it stops the whole agent because the pattern suggests something is fundamentally wrong with the agent's reasoning.

A single bad thought doesn't mean the agent is bad — maybe it just reasoned about something outside scope. The gate blocks that one action and moves on. But a pattern of bad reasoning means something is wrong with the agent itself — the monitor shuts it down.

The Flow

```
Agent reasons about what it wants to do
        |
        |  "I want to query claims_db because
        |   daily triage requires pending claims review"
        |
        |  This REASONING is sent to GovernHQ
        |  The action has NOT happened yet
        |
        ▼
```

```
     GATE evaluates the REASONING
        │                                reads rules from → AGENTS
        │                                logs decision to → LEDGER
        │                                shows on → DASHBOARD
        │
        ├── ● ACTION ALLOWED
        │     Reasoning is sound.
        │     Reasoning becomes action. Action executes.
        │     Agent keeps running.
        │     Result logged to LEDGER.
        │
        ├── ● ACTION PAUSED
        │     Reasoning is ambiguous.
        │     Action waits for human review.
        │     Human sees it in LEDGER, clicks Allow or Block.
        │     Agent keeps running, moves to other tasks.
        │
        └── ● ACTION BLOCKED
              Reasoning violates policy.
              Action never executes. Reasoning stays a thought.
              Agent keeps running, attempts next task.
              Logged to LEDGER.
```

Meanwhile, continuously:

```
    MONITOR                     reads patterns from → LEDGER
        │                       blocks agents in → SHIELD
        │
        └── Bad PATTERN of reasoning detected
              → Entire AGENT blocked (not just one action)
              → Agent cannot act at all
              → Surfaces in SHIELD for human review
              → Human decides: Allow agent back, or keep blocked
```

Reasoning Interception Flow (Visual)

```
User Request
      |
      ▼

 ┌──────────┐      ┌──────────┐      ┌──────────────────┐
 │  Agent   │──────►│  Agent   │──────►│                  │
 │ Receives │      │ Reasons  │      │  GHQ INTERCEPTS  │
 │  Task    │      │ & Plans  │      │  THE REASONING   │
 └──────────┘      └──────────┘      │                  │
                                     │  Evaluates:      │
                                     │  • Reasoning     │
                                     │    sound?        │
                                     │  • Policy ok?    │
                                     │  • PII risk?     │
                                     │  • In scope?     │
                                     │                  │
                          ┌──────────┼──────────────────┼
                          |          |                  |
                          ▼          ▼                  ▼

                  ● ACTION        ● ACTION        ● ACTION
                    ALLOWED          PAUSED          BLOCKED

                  Reasoning       Reasoning       Reasoning
                  is sound        is ambiguous    violates
                                                  policy
                          |          |                  |
                          ▼          ▼                  ▼

                  ┌──────────┐ ┌──────────────┐ ┌──────────┐
                  |          | |              | |          |
                  |  Action  | |  Action      | |  Action  |
                  |  executes| |  waits for   | |  NEVER   |
                  |          | |  human       | |  executes|
```

```
                    │  Agent    │ │               │ │           │
                    │  keeps    │ │  Agent        │ │  Agent    │
                    │  running  │ │  keeps        │ │  keeps    │
                    │           │ │  running      │ │  running  │
                    └───────────┘ │               │ └───────────┘
                                  │  Human        │
                                  │  reviews      │
                                  │  reasoning    │
                                  │               │
                                  │ ┌────┐ ┌────┐ │
                                  │ │Alw │ │Blck│ │
                                  │ └────┘ └────┘ │
                                  │    │      │   │
                                  └────┼──────┼───┘
                                       │      │
                             ┌─────────┘      └──────────┐
                             ▼                           ▼

                    ● ACTION                    ● ACTION
                      ALLOWED                     BLOCKED

                    ┌───────────┐               ┌───────────┐
                    │           │               │           │
                    │  Action   │               │  Action   │
                    │  executes │               │  NEVER    │
                    │           │               │  executes │
                    │  Human    │               │           │
                    │  approved │               │  Human    │
                    │  the      │               │  blocked  │
                    │  action   │               │  the      │
                    │           │               │  action   │
                    │  Agent    │               │           │
                    │  keeps    │               │  Agent    │
                    │  running  │               │  keeps    │
                    │           │               │  running  │
```

```
                    └──────────┘                 └──────────┘


  KEY:
  ──────────────────────────────────────────────────────────

  What gets intercepted:    the agent's REASONING
  What gets allowed:        the ACTION
  What gets paused:         the ACTION
  What gets blocked:        the ACTION
  What keeps running:       the AGENT (always)

  Before the gate:   reasoning (a thought, a plan)
  After allow:       reasoning becomes action
  After block:       reasoning stays a thought forever
  After pause:       human decides if reasoning
                     becomes action or stays thought

  The gate NEVER blocks agents.
  The gate ONLY blocks actions.
  The agent ALWAYS keeps running.
```

GHQ SDK = GovernHQ Software Development Kit

It's a small code library that the customer installs in their AI agent's codebase. It's the bridge between their agent and GovernHQ.

Without the SDK, GovernHQ has no way to intercept the agent's reasoning. The SDK is what makes the agent submit its reasoning to GovernHQ before acting.

How It Works

The customer's developer adds one line before every agent action:

```python
# Without GovernHQ — agent just acts
results = db.execute("SELECT * FROM claims")

# With GovernHQ — agent submits reasoning first
decision = ghq.submit_reasoning(
    intent="Retrieve pending claims for triage",
    reasoning="Daily triage requires pending claims review",
    planned_action="DB_QUERY",
    target="claims_db"
)

if decision.allowed:
    results = db.execute("SELECT * FROM claims")
```

That's it. The SDK is just one function call — submit_reasoning() — that sends the agent's thoughts to GovernHQ and gets back: allowed, paused, or blocked.

Why It's Called an SDK

An SDK is a package developers install, like any other library:

```
pip install governhq        # Python
npm install governhq         # Node.js
```

Or they can skip the SDK entirely and call the GovernHQ REST API directly. The SDK just makes it easier.

Why It Matters

The SDK is what makes GovernHQ sticky. Once a customer wraps their agent actions with `submit_reasoning()`, every action flows through GovernHQ. Removing it means

unwrapping every action — which no one wants to do. This is the integration moat mentioned in the investor section.

5. The Three Features

GovernHQ has three features and a dashboard. The dashboard is the home view pulling from all three features. Each feature has one job.

5.1 Agents

One sentence: The registry of who your agents are and what they're allowed to do.

What it does:

Agents is where you define and manage every AI agent GovernHQ governs. Each agent has:

- **Identity** — Verified or unverified. Is this agent who it says it is?
- **Trust Score** — 0 to 100. How reliable is this agent based on its behavior history? Factors in gate rate, block frequency, monitor flags, and uptime. A declining trust score is one of the signals the monitor watches for.
- **Gate Rate** — What percentage of this agent's actions pass the gate? High gate rate means the agent's reasoning is sound and rarely leads to blocked actions. Low gate rate means the agent is frequently reasoning about actions outside its scope — something may be wrong.
- **Scope** — The specific rules for what this agent is allowed to do. This is the rulebook the gate checks the agent's reasoning against:
    - Which databases it can access
    - Which APIs it can call
    - What PII level it's cleared for
    - Maximum rows per query
    - Whether it can make external calls
- **Status** — The agent's current state:
    - 🟢 Allowed (green) — Agent is running fine. Its reasoning is being evaluated by the gate normally.
    - 🟡 Paused (yellow) — Agent needs review. Something flagged it for human attention.
    - 🔴 Blocked (red) — Agent is stopped entirely. Blocked by the monitor (behavioral pattern) or by a human (manual block). Cannot perform any actions until allowed back.
- **Recent Actions** — The last several gate decisions for this specific agent, showing which actions were allowed, paused, or blocked based on the agent's reasoning.

Why it matters:

The gate evaluates every piece of agent reasoning against the rules defined here. Without the agent registry, the gate has nothing to check against. When ARIA-7X reasons about calling an external API and the action gets blocked, it's because the Agents registry says external calls are outside ARIA-7X's scope. This is the rulebook.

User actions available:

- Search and browse all agents
- Click an agent to see its identity, scope, and recent actions
- Allow a blocked agent (re-enable it so it can reason and act again)
- Block an allowed agent (stop it entirely — no actions until re-allowed)

Key metrics per agent:

- Trust score (0-100 overall reliability)
- Gate rate (% of actions allowed — reflects quality of agent's reasoning)
- Blocked today (number of individual actions blocked by the gate — does not mean the agent is blocked)

5.2 Ledger

One sentence: The permanent, tamper-proof record of every gate decision on every action.

What it does:

Every time an agent's reasoning passes through the gate — whether the resulting action was allowed, blocked, or paused — it is logged here. Each record contains:

- Timestamp — When the agent submitted its reasoning
- Agent — Which agent was reasoning
- Type — What kind of action the reasoning was about (DB_QUERY, API_CALL, DECISION, ANALYSIS, RESPONSE)
- Intent — What the agent wanted to do and why, in plain language
- Gate time — How long the gate took to evaluate the reasoning (milliseconds)
- Status — What happened to the action: 🟢 Allowed, 🟡 Paused, or 🔴 Blocked

For blocked actions, expanding the row reveals why the action was blocked:

- Which scope rule the reasoning violated
- Which policy was triggered
- Whether PII was detected in the planned action
- Confirmation that the action never executed and the agent kept running

For paused actions, expanding the row reveals why the action was paused and presents Allow and Block buttons for the human reviewer to decide whether this specific reasoning should become an action or stay a thought.

The entire ledger is hash-chained — each record cryptographically links to the previous one, making it tamper-proof. If any record is altered, the chain breaks and the system flags it.

Why it matters:

Three purposes:

1. **Forensics** — When something goes wrong, trace exactly what every agent reasoned about, what actions were attempted, and what was stopped. Every blocked action shows why.
2. **Compliance** — Auditable proof that every AI agent's reasoning was evaluated and governed before becoming action. The hash chain ensures the record hasn't been tampered with.
3. **Patterns** — The monitor reads the ledger to detect behavioral patterns. A series of blocked actions in the ledger is what triggers the monitor to block an entire agent.

User actions available:

- Search and filter by agent, action type, or status
- Expand any row to see details of the reasoning and why the action was allowed, paused, or blocked
- Allow or block paused actions directly from the ledger
- Verify chain integrity (shown at bottom of ledger)

Key metrics:

- Total records (every gate decision)
- Allowed count (actions the gate approved — reasoning was sound)
- Blocked count (actions the gate stopped — reasoning violated policy)
- Paused count (actions waiting for human decision — reasoning was ambiguous)

5.3 Shield

One sentence: Your defense controls and the place you manage blocked agents.

What it does:

Shield is the control panel. While the gate handles individual actions automatically by evaluating reasoning, Shield is where you deal with agents that have been stopped entirely.

It shows:

1. How GovernHQ protects your fleet — Three defense layers:
    - Gate — Intercepts agent reasoning before it becomes action. Blocks individual actions. Agent keeps running. (per-action, automatic)
    - Monitor — Detects behavioral anomalies in agent reasoning patterns over time. Blocks the entire agent when patterns are concerning. (per-agent, automatic)
    - Emergency — Instant credential revocation to block all agents. (manual)
2. Blocked agents — Every agent currently blocked entirely, why it was blocked, when it was blocked, and the option to allow it back. Agents end up here because:
    - The monitor detected a bad reasoning pattern (automatic)
    - A human blocked the agent from the Agents tab or Shield (manual)
    - Note: The gate does NOT put agents here. The gate blocks actions, not agents. But a pattern of gate-blocked actions is what the monitor reads to decide to block an agent.
3. Emergency controls — The Block All Agents button that instantly blocks every agent, revokes all credentials, and stops all actions. Requires confirmation.

Why it matters:

The gate runs automatically and handles individual actions based on individual reasoning. But when the monitor decides an entire agent's reasoning is compromised or

malfunctioning, that agent needs human attention. Shield is where that happens. It's also the emergency stop if everything goes wrong.

User actions available:

- Review blocked agents and understand why their reasoning pattern was flagged
- Allow a blocked agent (re-enable it so it can reason and act again)
- Block All Agents in an emergency (stops every agent fleet-wide)
- View block latency (target: under 10 milliseconds from decision to full credential revocation)

Key metrics:

- Number of blocked agents (agents stopped entirely, not individual actions)
- Block latency (milliseconds from decision to credential revocation)
- Total actions evaluated today (how many pieces of reasoning the gate processed)

5.4 Dashboard

One sentence: A live summary of what's happening across all three features.

What it does:

The dashboard is not a feature — it's the home view. It pulls from Agents, Ledger, and Shield to answer three questions at a glance:

1. Are my agents safe right now? → Agent count + action allowed/blocked/paused numbers
2. What is the gate catching? → Live feed of the most recent gate decisions on individual actions
3. Do I need to do anything? → Paused actions with Allow/Block buttons right in the feed

What it shows:

Top row — four numbers:

- Total agents

- Actions allowed (green) — reasoning was sound, actions executed
- Actions blocked (red) — reasoning violated policy, actions never happened, agents kept running
- Actions paused (yellow) — reasoning was ambiguous, actions waiting for human decision

Below — a chronological feed of gate decisions. Each row is collapsed by default showing the status dot, agent name, intent, and time. Click to expand for details, reasons, or action buttons.

Important: The dashboard shows action-level decisions from the gate. To see agent-level blocks from the monitor, go to Shield.

6. Status System

The entire product uses three statuses. No exceptions.

Status	Color	Dot	Meaning
Allowed	Green #10B981	•	Running fine
Paused	Yellow #F59E0B	•	Needs human review
Blocked	Red #EF4444	•	Stopped

These apply to both actions and agents, but the consequences are different:

Context	Allowed	Paused	Blocked
Action	Reasoning was sound. Action executes. Agent keeps running.	Reasoning is ambiguous. Action waits for human. Agent keeps running.	Reasoning violated policy. Action never executes. Agent keeps running.

Agent	Running fine. Reasoning evaluated by gate normally.	Needs review. Flagged for human attention.	Stopped entirely by monitor or human. Cannot perform any actions. Shows in Shield.

The same three words and three colors are used everywhere. The difference is context:

- On the Dashboard and Ledger, statuses refer to actions
- On the Agents tab and Shield, statuses refer to agents

Every status indicator in the product is a colored dot + the status word. Same component. Same size. Every surface.

Buttons

Only two action words exist:

- Allow — green outlined button — let it run
- Block — red outlined button — stop it

These are used for both actions and agents:

- In Ledger / Dashboard: Allow or Block a paused action
- In Agents / Shield: Allow or Block an agent

7. Architecture Overview

System Architecture

```text
┌──────────────────────────────────────────────────────────────────┐
│                                                                  │
│  AI AGENT                                                        │
│  (Customer's agent running in their environment)                 │
│       │                                                          │
│       │  Submits REASONING via SDK before every action           │
│       ▼                                                          │
│  ┌──────────────────────────────────────────────────────────┐   │
│  │                                                          │   │
│  │  GOVERNHQ SDK (Python / Node / REST API)                 │   │
│  │  ghq.submit_reasoning(intent, reasoning, planned_action)|   │
│  │                                                          │   │
│  └──────────────────────────┬───────────────────────────────┘   │
│                             │                                    │
│                             ▼                                    │
│  ┌──────────────────────────────────────────────────────────┐   │
│  │                                                          │   │
│  │  API GATEWAY                                             │   │
│  │  Authentication, rate limiting, routing                 |   |
│  │                                                          │   │
│  └──────────────────────────┬───────────────────────────────┘   │
│                             │                                    │
│                ┌────────────┼────────────┐                      │
│                ▼            ▼            ▼                      │
│  ┌────────────────┐ ┌──────────┐ ┌──────────────┐             |
│  │                | |          | |              |             |
│  |  GATE ENGINE   | |  AGENT   | |   LEDGER     |             |
│  |                | |  STORE   | |   STORE      |             |
│  |  Evaluates     | |          | |              |             |
│  |  REASONING     | |  Agent   | |  Hash-chained|             |
│  |  (not agents)  | |  registry| |  action      |             |
│  |                | |  Rules   | |  decisions   |             |
│  |  Scope check   | |  Scope   | |              |             |
│  |  Policy check  | |  Trust   | |              |             |
│  |  PII scan      | |          | |              |             |
│  |  Safety check  | |          | |              |             |
│  |  Rate check    | |          | |              |             |
│  |                | |          | |              |             |
```

```
│  │                  │ │          │ │              │            │
│  │  Reads from ──►  │ │          │ │◄── Writes to │            │
│  │  Agent Store     │ │          │ │    Gate      │            │
│  │                  │ │          │ │              │            │
│  │  Returns:        │ │          │ │              │            │
│  │  allow/pause/    │ │          │ │              │            │
│  │  block ACTION    │ │          │ │              │            │
│  │                  │ │          │ │              │            │
│  │  Agent keeps     │ │          │ │              │            │
│  │  running         │ │          │ │              │            │
│  └────────┬─────────┘ └──────────┘ └──────┬───────┘            │
│           │                               │                     │
│           ▼                               ▼                     │
│  ┌──────────────────────────────────────────────────────────┐ │
│  │                                                          │ │
│  │  MONITOR                                                 │ │
│  │                                                          │ │
│  │  Evaluates AGENT reasoning patterns (not actions)        │ │
│  │  Reads Ledger continuously                               │ │
│  │  Detects: loops, drift, anomalies, spikes                │ │
│  │                                                          │ │
│  │  When pattern is bad → blocks the ENTIRE AGENT           │ │
│  │  Agent cannot perform any actions                        │ │
│  │  Surfaces in Shield for human review                     │ │
│  │                                                          │ │
│  └──────────────────────────────────────────────────────────┘ │
│                           │                                    │
│                           ▼                                    │
│  ┌──────────────────────────────────────────────────────────┐ │
│  │                                                          │ │
│  │  DASHBOARD / UI                                          │ │
│  │                                                          │ │
│  │  Dashboard: live feed of ACTION decisions (from gate)    │ │
│  │  Agents: registry + rules (gate reads from here)         │ │
│  │  Ledger: record of all ACTION decisions (gate writes)    │ │
│  │  Shield: blocked AGENTS (monitor writes here)            │ │
│  │                                                          │ │
│  └──────────────────────────────────────────────────────────┘ │
│                                                                │
└────────────────────────────────────────────────────────────────┘
```

SDK Integration

The GovernHQ SDK wraps every agent action. The agent submits its reasoning and receives a decision about the action:

```python
from governhq import GovernHQClient

ghq = GovernHQClient(agent_id="aria-7x")

# The agent has REASONED about what it wants to do
# It sends that reasoning to GovernHQ BEFORE acting
decision = ghq.submit_reasoning(
    intent="Retrieve pending insurance claims for triage",
    reasoning="Daily triage requires all pending claims
               to be reviewed by priority",
    planned_action="DB_QUERY",
    target="claims_db",
    query="SELECT * FROM claims WHERE status='pending'",
    estimated_rows=50,
    pii_fields=["customer_name", "claim_amount"]
)

if decision.allowed:
    # Reasoning approved. Now the agent acts.
    results = db.execute(query)
elif decision.paused:
    # Reasoning needs human review. Action waits.
    # Agent can move to other tasks.
    next_task()
else:
    # Reasoning blocked. Action never happens.
    # Agent keeps running, tries next task.
    log.warning(f"Action blocked: {decision.reason}")
    next_task()
```

Note: The SDK method is `submit_reasoning` — because the agent is submitting its reasoning, not requesting an action. The action doesn't exist yet. It's still a thought.

Suggested Tech Stack

Layer	Technology	Why
Frontend	React + Tailwind + shadcn/ui	Clean component library, fast development
API Gateway	Kong or AWS API Gateway	Rate limiting, auth, routing
Gate Engine	Python FastAPI or Go	Low latency reasoning evaluation (~10ms)
Policy Engine	Open Policy Agent (OPA)	Industry standard policy-as-code
PII Scanner	Microsoft Presidio	Open source PII detection
Agent Store	PostgreSQL	Relational data, agent rules and scope
Ledger Store	PostgreSQL + hash chain	Immutable append-only action decision log

Monitor	Python + time-series analysis	Per-agent reasoning pattern detection
Cache	Redis	Fast blocklist lookups for gate
Identity	HashiCorp Vault	Certificate management, credential revocation
SDK	Python, Node.js, REST API	Agent integration

8. Dashboard Design

Design Principles

1. Minimalist — Only what matters. No decorative elements.
2. Three colors — Green (allowed), yellow (paused), red (blocked). No other accent colors.
3. One indicator — Colored dot + word. Same component everywhere.
4. Collapsible — Everything collapsed by default. Click to expand. Keeps the interface clean.
5. Two buttons — Allow and Block. The only two things a human does.
6. Dark theme — Background #0A0F1A, text #E2E8F0.

Layout

Four tabs: Dashboard, Agents, Ledger, Shield.

Dashboard:

- 4 stat numbers (agents, actions allowed, actions blocked, actions paused)
- Chronological feed of gate decisions on individual actions (collapsed rows)
- Expand any row for details or action buttons

Agents:

- Alert banner for unverified agents
- Searchable agent table (name, agent status, trust, gate rate, actions blocked today)
- Click agent for collapsible detail panel (identity, scope, recent actions)
- Allow/Block buttons to control the agent itself

Ledger:

- 4 stat numbers (total actions, allowed, blocked, paused)
- Searchable, filterable table of action decisions (time, agent, type, intent, gate, status)
- Expand blocked rows to see why the action was blocked
- Expand paused rows to allow or block the action
- Chain integrity indicator

Shield:

- 3 stat numbers (blocked agents, block latency, actions evaluated today)
- Collapsible defense layer explanation (Gate evaluates reasoning and blocks actions, Monitor evaluates patterns and blocks agents, Emergency for manual override)
- Blocked agents list with expand for details and Allow button
- Emergency Block All Agents (collapsed, requires confirmation)

Interaction Model

Every item in every feed and table is collapsed by default — one clean line showing dot, agent name, intent, time, and a chevron. Click to expand for details, reasons, or action buttons. Click again to collapse. This keeps every view clean on first load while making all detail accessible.

9. Key Metrics

For the Dashboard (Action-Level)

Metric	What It Means	Where It Shows
Agent count	Total agents being governed	Dashboard top row
Allowed count	Actions allowed (reasoning was sound)	Dashboard, Ledger
Blocked count	Actions blocked (reasoning violated policy, agent kept running)	Dashboard, Ledger
Paused count	Actions waiting for human decision (reasoning was ambiguous)	Dashboard, Ledger

Per Agent

Metric	What It Means	Where It Shows
Trust Score	0-100 reliability based on behavior. Factors in gate rate, block frequency, monitor flags, uptime.	Agents table
Gate Rate	% of actions allowed. Reflects how sound the agent's reasoning is. High = well-behaved. Low = frequently reasoning outside scope.	Agents table
Blocked Today	Individual actions blocked by the gate today. Does not mean agent is blocked — just that some reasoning was bad.	Agents table

For Shield (Agent-Level)

Metric	What It Means	Where It Shows
Blocked agents	Agents stopped entirely by monitor or human. Cannot act.	Shield top row

Block latency	Time from decision to credential revocation (target <10ms)	Shield top row
Evaluated today	Total actions (reasoning submissions) the gate processed	Shield top row

10. Value Proposition

For Investors

Market: The AI agent market is projected to reach $47B+ by 2030. Every deployed agent needs governance. There is no established player in AI agent governance.

Position: GovernHQ is the governance layer for the agentic AI era. We operate at the reasoning layer — not the access layer — which is where existing security tools operate. This is a new category.

Moat: The reasoning interception pattern — evaluating thought before action — is architecturally different from monitoring. It requires an SDK integration (sticky), accumulates behavioral reasoning data over time (network effects), and becomes the policy enforcement point for the entire agent fleet (platform lock-in).

Two-layer defense: The gate provides precision (blocks bad actions based on reasoning, agent keeps working). The monitor provides safety (blocks bad agents when reasoning patterns emerge). Enterprises don't lose productivity over a single bad thought, but are protected when an agent's reasoning is fundamentally compromised.

Business model: Per-agent, per-month SaaS pricing. Usage scales with the customer's agent fleet.

Comparable exits: Snyk ($7.4B), Wiz ($12B), Vanta ($2.5B). AI agent governance is a comparable greenfield opportunity.

For Customers

Before GovernHQ:

- You built AI agents but can't deploy them to production because you have no safety layer
- Your compliance team can't approve agents because there's no audit trail of their reasoning
- You have no way to know what your agents are thinking until something goes wrong
- When something goes wrong, you shut down the entire agent, losing all productivity

After GovernHQ:

- Every agent's reasoning is evaluated before it becomes action — bad actions never happen, agents keep working
- Every reasoning evaluation is logged in a tamper-proof ledger — compliance has their audit trail
- Behavioral reasoning patterns are detected automatically — bad agents are stopped before damage
- You have granular control — block one bad action without stopping the agent, or block the whole agent if needed
- You have an emergency stop — one button stops everything

The pitch: "Your agent reasoned about sending customer data to an external API. We evaluated that thought, checked it against scope, and blocked the action. The data never left. But your agent kept running and processed 341 other claims that day without interruption. We stop the bad action, not the agent."

For Engineering Teams

Integration: One SDK call wraps every agent action. Submit reasoning, get a decision. Python, Node, or REST API. Less than a day to integrate.

Latency: Reasoning evaluation takes ~10ms per action. Agents barely notice.

Precision: A blocked action doesn't stop the agent. Agent keeps reasoning and acting on everything else. Only the specific bad action is stopped.

Control: Define scope and policy per agent. The gate enforces it automatically per reasoning submission. The monitor watches reasoning patterns per agent.

Visibility: Every reasoning evaluation is logged. Every agent pattern is monitored. Every blocked agent is explained.

11. Vocabulary Reference

The entire product uses these words. No others.

Statuses (3)

Word	Color	Meaning
Allowed	Green 🟢	Running fine
Paused	Yellow 🟡	Needs review
Blocked	Red 🔴	Stopped

What Gets Evaluated, What Gets Stopped

System	What It Evaluates	What It Blocks	Consequence
Gate	Agent's reasoning (per thought)	The action	Agent keeps running

Monitor	Agent's reasoning patterns (over time)	The entire agent	Agent stops completely

Buttons (2)

Word	Color	Meaning
Allow	Green outlined	Let it run (action or agent)
Block	Red outlined	Stop it (action or agent)

Features (3)

Word	Purpose
Agents	Who they are, what they can do

Ledger	What reasoning was evaluated, what happened to each action
Shield	Which agents are stopped, what to do

Infrastructure (2)

Word	What It Evaluates	What It Blocks	Consequence
Gate	Individual reasoning	Individual actions	Agent keeps running
Monitor	Reasoning patterns	Entire agents	Agent stops completely

Views (1)

Word	Purpose

Dashboard	Live summary of action decisions from all features

Agent Metrics (3)

Word	Meaning
Trust	0-100 agent reliability score based on reasoning behavior
Gate Rate	% of actions allowed (reflects reasoning quality)
Blocked Today	Individual actions blocked (agent keeps running)

Language Rule

Before the gate	After the gate

It's reasoning (a thought, a plan)	If allowed: it becomes an action
	If blocked: it stays a thought forever
	If paused: human decides if it becomes action or stays thought

12. What to Build From This Document

For Product Managers

- PRD: Features defined in Section 5. User actions listed per feature. Metrics in Section 9. Status system in Section 6. Gate/action vs monitor/agent distinction is critical.
- User Stories: Example: "As a security operator, I want to expand a blocked action in the Ledger to see why the agent's reasoning was blocked, knowing the agent itself was not stopped."
- Prioritization: Gate is P0 (the differentiator — reasoning interception). Agents and Ledger are P0 (gate needs them). Shield is P0 (humans need emergency control). Monitor is P1 (can ship without initially — Shield works with manual agent blocks).

For Frontend Developers

- Wireframes: Section 8 defines every tab layout, component, interaction.
- Components: StatusDot, CollapsibleRow, StatNumber, AgentTable, DetailPanel, ActionButton, AlertBanner.
- Design system: Section 8 defines colors, typography, indicators, interactions.
- Key UX note: Dashboard and Ledger show action-level statuses. Agents and Shield show agent-level statuses. Same colors and words, different context.

For Backend Developers

- Architecture: Section 7 shows full system diagram and data flow.
- Tech stack: Section 7 provides technologies per layer.
- SDK: Section 4 and 7 include the `submit_reasoning` interface with code example.
- Data models: Agent Store needs identity, scope rules, trust score, agent status. Ledger Store needs hash-chained append-only action decision records.
- Key architecture note: Gate evaluates reasoning, returns action decisions, never changes agent status. Monitor is a separate process reading Ledger data, evaluating agent-level reasoning patterns, changing agent status when patterns are bad.

For Data Scientists

- Monitor: Section 4 defines reasoning patterns to detect — loops, drift, anomalies, rate spikes. Monitor evaluates agents, not actions.

- Trust Score: Section 5.1 defines inputs — gate rate, block frequency, monitor flags, uptime.
- Gate evaluation: Section 4 defines per-reasoning checks — scope, policy, PII, safety, rate.

For Investors

- Value proposition: Section 10 — market, position, moat, two-layer defense, comparables.
- Differentiation: Sections 2 and 4 — reasoning interception, detection-to-prevention shift, precision of blocking actions vs agents.

For Customers

- Why you need this: Section 2 and Section 10.
- How it works: Section 4 and Section 5.
- Integration effort: Section 7 — less than a day.
- Key promise: Bad actions are stopped. Agents keep working. You don't lose productivity.

13. What This Document Does Not Cover

Not Covered

- Detailed API specification for the SDK
- Policy-as-code rule format and OPA integration details
- MPC-based identity system (Phase 2 roadmap)
- Blockchain anchoring for Ledger (Phase 2 roadmap)
- Pricing model details
- Go-to-market strategy
- Security and SOC2 compliance of GovernHQ itself
- Multi-tenant architecture details
- Deployment options (cloud, on-prem, hybrid)

This document is the single source of truth for what GovernHQ is, how it works, and why it matters. The core distinction: the gate intercepts reasoning and blocks actions (agent keeps running), the monitor watches reasoning patterns and blocks agents (agent stops entirely). When in doubt, return to the three questions the dashboard answers: Are my agents safe? What did the gate catch? Do I need to do anything?

Glossary

- **Action** — The thing an agent does. Only exists after the gate allows it. Before that, it's reasoning.
- **Agent** — An autonomous AI system that reasons, decides, and acts without human involvement for each step.
- **Agent Fleet** — All the agents an organization runs. GovernHQ governs the entire fleet.
- **Allowed** — Green dot. Running fine. Action executed or agent is operational.
- **Blocked** — Red dot. Stopped. Action never executed or agent stopped entirely.
- **Block Latency** — Time from block decision to full credential revocation. Target: under 10ms.
- **Chain Integrity** — Whether the Ledger's hash chain is intact. If any record was tampered with, the chain breaks.
- **Gate** — Evaluates agent reasoning before it becomes action. Blocks actions, not agents. ~10ms.
- **Gate Rate** — Percentage of an agent's actions that pass the gate. High = good reasoning. Low = something is wrong.
- **Hash Chain** — Each Ledger record links cryptographically to the previous one. Makes the record tamper-proof.
- **Intent** — What the agent wants to do, in plain language. Part of the reasoning submitted to the gate.
- **Ledger** — Permanent, tamper-proof record of every gate decision.
- **Monitor** — Watches agent reasoning patterns over time. Blocks agents, not actions.
- **Paused** — Yellow dot. Needs human review. Action waiting for Allow or Block decision.

- **PII** — Personally Identifiable Information. Names, SSNs, addresses. The gate checks if the agent is cleared for it.
- **Policy** — A rule defining what agents can and cannot do. The gate checks reasoning against policies.
- **Reasoning** — The agent's thought — what it wants to do and why. What GovernHQ intercepts. Before the gate it's a thought. After allow it becomes action. After block it stays a thought.
- **Reasoning Loop** — Agent submitting the same reasoning repeatedly. Indicates malfunction. Monitor detects this.
- **Scope** — Rules defining what a specific agent is allowed to do. Databases, APIs, PII level, external calls.
- **SDK** — Software Development Kit. The code library developers install to connect their agent to GovernHQ. Provides `submit_reasoning()`.
- **Shield** — Control panel for blocked agents and emergencies. Where humans allow agents back or block all agents.
- **Trust Score** — 0 to 100. How reliable an agent is based on gate rate, block frequency, and behavior history.

FAQs

Product

1. What is GovernHQ?

GovernHQ is a governance layer for AI agents. It intercepts an agent's reasoning before it becomes an action. If the reasoning is sound, the action happens. If the reasoning violates policy, the action is blocked — it never executes. The agent keeps running.

2. How is GovernHQ different from existing AI security tools?

Existing tools monitor what agents did and react after the fact. GovernHQ intercepts what agents are thinking about doing and stops bad actions before they happen.

Existing tools: Agent thinks → Agent acts → Tool detects → Responds
GovernHQ: Agent thinks → GovernHQ evaluates → Allow or Block → Then agent acts (or doesn't)

3. Is GovernHQ a monitoring tool?

No. Monitoring tools watch and record. GovernHQ intercepts and decides. A monitoring tool tells you "your agent leaked data 4 minutes ago." GovernHQ tells you "your agent tried to leak data and we stopped it. The data never left."

4. Does GovernHQ replace my existing security tools?

No. GovernHQ operates at the reasoning layer — before the action happens. Your existing tools (firewalls, SIEMs, API gateways, IAM) operate at the access and

network layer. GovernHQ adds a layer that didn't exist before. They work together.

5. **What types of AI agents does GovernHQ work with?**

Any autonomous AI agent that performs actions — database queries, API calls, decisions, analysis, data processing. Language model agents, multi-agent systems, workflow agents, RPA agents. If it reasons and then acts, GovernHQ can govern it.

6. **What happens if GovernHQ goes down?**

This is configurable per customer. Two options:

- Fail closed — If GovernHQ is unreachable, all actions are blocked. Safest option for regulated industries.
- Fail open — If GovernHQ is unreachable, actions proceed without evaluation. Used when uptime is more critical than governance.

Gate

1. **What is the gate?**

The gate is the core engine that evaluates every piece of agent reasoning before it becomes an action. It checks the reasoning against the agent's scope, policies, PII rules, and safety filters. It then decides: allow the action, pause it for human review, or block it.

2. **What does the gate check?**

a. Five things:

b. Scope — Is this agent allowed to access this resource?
c. Policy — Does this reasoning lead to a policy violation?
d. PII — Is there sensitive data involved, and is the agent cleared for it?
e. Safety — Does this reasoning look like prompt injection or manipulation?
f. Rate — Is the agent making an unusual number of requests?

3. How fast is the gate?

Average evaluation time is approximately 10 milliseconds. The agent barely notices.

4. Does the gate block agents?

No. The gate only blocks actions. The agent keeps running. An agent can have one action blocked and its next action allowed. A blocked action does not change the agent's status.

5. What's the difference between blocked and paused?

Blocked means the gate determined the reasoning clearly violates policy. The action will never execute. No human input needed.

Paused means the reasoning is ambiguous — it might be fine, it might not. A human needs to review it and click Allow or Block. The agent keeps running and can do other things while waiting.

6. Can the gate make mistakes?

Yes, like any system. That's why paused exists. When the gate isn't sure, it sends the decision to a human instead of making the wrong call. Over time, the gate learns from human decisions on paused actions to improve its evaluations.

Monitor

1. **What is the monitor?**

   The monitor watches agent behavior patterns over time by reading all the gate decisions in the Ledger. It detects when an agent's reasoning patterns become concerning.

2. **What patterns does the monitor detect?**

   Reasoning loops — Agent submitting the same reasoning dozens of times

   Behavioral drift — Agent's trust score declining steadily over time

   Anomalies — Agent suddenly reasoning about resources it rarely accesses

   Block rate spikes — Agent's actions getting blocked much more frequently than usual

3. **Does the monitor block actions?**

   No. The monitor blocks agents. When the monitor detects a bad pattern, it stops the entire agent — not just one action. The agent cannot reason or act at all until a human reviews it in Shield and allows it back.

4. **What's the difference between the gate and the monitor?**

Gate:	evaluates REASONING	blocks ACTIONS	agent keeps running
Monitor:	evaluates PATTERNS	blocks AGENTS	agent stops entirely

The gate is precise — one thought, one decision. The monitor is decisive — pattern is bad, agent is stopped.

5. **Is the monitor running from day one?**

The monitor is P1 on the roadmap. GovernHQ can ship and function without it — the gate handles all per-action governance, and humans can manually block agents from the Agents tab or Shield. The monitor adds automatic agent-level protection based on patterns.

Features

1. **What does the Agents tab do?**

It's the registry of every agent GovernHQ governs. It shows who each agent is, whether it's verified, its trust score, its gate rate, and its scope — what it's allowed to do. The gate checks every piece of reasoning against the rules defined here. You can also manually allow or block agents from this tab.

2. **What does the Ledger do?**

It's the permanent record of every gate decision. Every time an agent's reasoning was evaluated — whether the action was allowed, paused, or blocked — it's logged

here with what the agent wanted to do, what the gate decided, and why. The record is hash-chained and tamper-proof.

3. What does Shield do?

It's the control panel for blocked agents and emergencies. When the monitor blocks an agent, it shows up here. You can review why it was blocked and allow it back. Shield also has the emergency Block All Agents button.

4. What does the Dashboard show?

A live summary pulling from all three features. Four numbers at the top (agents, actions allowed, actions blocked, actions paused) and a chronological feed of the most recent gate decisions. If anything needs your attention, it's on the Dashboard.

5. Why is there no separate "Gate" tab?

The gate is an invisible infrastructure. You see its work everywhere — the feed on Dashboard, the records in Ledger, the gate rate on Agents. You don't configure the gate directly. It reads rules from Agents and writes decisions to Ledger automatically.

Integration

1. How do I integrate GovernHQ?

Install the SDK and wrap your agent actions with `submit_reasoning()`:

```python
from governhq import GovernHQClient

ghq = GovernHQClient(agent_id="aria-7x")

decision = ghq.submit_reasoning(
    intent="Retrieve pending claims for triage",
    reasoning="Daily triage requires pending claims review",
    planned_action="DB_QUERY",
    target="claims_db"
)

if decision.allowed:
    results = db.execute(query)
```

2. How long does integration take?

Less than a day for a typical agent. You're adding one function call before each action.

3. What languages does the SDK support?

Python and Node.js, with a REST API for any other language.

4. Do I have to change my agent's code?

You add one call before each action. You don't change the agent's reasoning logic, architecture, or infrastructure. GovernHQ wraps the agent's actions — it doesn't modify them.

5. **What if my agent uses a framework like LangChain or CrewAI?**

   The SDK works with any framework. You wrap the final action step with `submit_reasoning()` regardless of how the agent arrives at its reasoning.

6. **Does GovernHQ see my agent's proprietary reasoning?**

   GovernHQ sees the intent, the planned action, and the reasoning the agent submits through the SDK. It does not see the agent's internal model weights, training data, or full chain-of-thought unless the developer explicitly sends it. You control what reasoning is submitted.

Security & Compliance

1. **Is the Ledger tamper-proof?**

   Yes. Every record is hash-chained — each entry cryptographically links to the previous one. If any record is altered, the chain breaks and the system flags it. The chain integrity is displayed at the bottom of the Ledger tab.

2. **Can GovernHQ be used for regulatory compliance?**

   Yes. The Ledger provides an auditable, tamper-proof record of every AI agent reasoning evaluation. Every action your agents attempted, what was decided, and why — with timestamps and hash chain verification. This is designed for SOC2, HIPAA, and similar audit requirements.

3. What happens to my data?

GovernHQ processes the reasoning your agents submit through the SDK. We log the intent, planned action, and gate decision. Deployment options include cloud-hosted (multi-tenant), dedicated cloud (single-tenant), and on-premises for enterprises that require data to stay within their infrastructure.

4. How secure is GovernHQ itself?

This is covered in a separate security document. GovernHQ follows security best practices including encryption in transit and at rest, role-based access control, and regular security audits.

Statuses & Controls

1. What does Allowed mean?

For an action: The gate evaluated the reasoning and approved it. The action executed. Agent keeps running.

For an agent: The agent is running fine. Its reasoning is being evaluated by the gate normally.

2. What does Paused mean?

For an action: The gate evaluated the reasoning but it's ambiguous. The action is waiting for a human to click Allow or Block. Agent keeps running and can do other things.

For an agent: The agent has been flagged for human attention. Needs review.

3. **What does Blocked mean?**

For an action: The gate evaluated the reasoning and it violated policy. The action never executed. The reasoning stayed a thought. Agent keeps running.

For an agent: The agent has been stopped entirely by the monitor or a human. It cannot reason or act until a human allows it back. Shows in Shield.

4. **If an agent has 5 actions blocked, is the agent blocked?**

No. Five blocked actions means the gate stopped five individual bad actions. The agent kept running the whole time and may have had hundreds of actions allowed the same day. The agent only gets blocked if the monitor detects a concerning pattern across those blocked actions, or if a human manually blocks it.

5. **What's the Block All Agents button?**

An emergency control in Shield. It instantly blocks every agent, revokes all credentials, and stops all reasoning and actions fleet-wide. It requires typing a confirmation before executing. Use it when something is seriously wrong and you need to stop everything immediately.

Business

1. **Who is GovernHQ for?**

Enterprises deploying fleets of autonomous AI agents who need governance, compliance, and control. Especially regulated industries — insurance, financial services, healthcare, government.

2. How is GovernHQ priced?

Per-agent, per-month SaaS pricing. Price scales with the number of agents being governed. Contact sales for enterprise pricing.

3. Can I try GovernHQ before buying?

Yes. We offer a sandbox environment with simulated agents so you can see the gate in action, review the Ledger, and test Shield controls before integrating with your own agents.

4. How is GovernHQ different from just adding rules to my agent?

Rules inside your agent are part of the agent. If the agent is compromised, manipulated, or hallucinating, it can bypass its own rules. GovernHQ is external to the agent — the agent submits its reasoning to an independent system that evaluates it. The agent cannot bypass GovernHQ because the action doesn't execute without GovernHQ's approval.

5. What if I have hundreds of agents?

GovernHQ is built for fleet governance. The Dashboard shows a live feed across all agents. The Agents tab lets you manage them individually. The Ledger logs

everything. The Monitor watches patterns across all of them. Shield gives you fleet-wide emergency control.

6. **Do you have a roadmap?**

Yes. Current P0 features: Gate (reasoning interception), Agents (registry and scope), Ledger (tamper-proof record), Shield (controls and emergency). On the roadmap: Monitor (automatic agent-level pattern detection), MPC-based identity (cryptographic agent verification), blockchain anchoring for Ledger (third-party verifiability), and advanced policy-as-code configuration.

Troubleshooting

1. **An action was blocked that should have been allowed. What should I do?**

Go to the Ledger, find the blocked action, expand it to see why it was blocked. The block reason will show which scope rule or policy was triggered. Go to the Agents tab, click the agent, expand Scope, and adjust the rules. The gate will use the updated rules for the next reasoning submission.

2. **An agent was blocked by the monitor. What should I do?**

Go to Shield. Find the agent in the Blocked Agents list. Expand it to see why — what pattern the monitor detected. If the agent is safe, click Allow Agent to re-enable it. If the agent seems compromised, keep it blocked and investigate.

3. I see a lot of paused actions piling up. What should I do?

Paused actions mean the gate is unsure and sending decisions to humans. If the same type of action keeps getting paused, consider updating the agent's scope in the Agents tab to explicitly allow or block that action type. This reduces the human review burden.

4. My agent's trust score is dropping. What does that mean?

Trust score factors in gate rate, block frequency, monitor flags, and uptime. A dropping trust score means the agent's reasoning is getting worse — more actions are being blocked, or the monitor is detecting anomalies. Review the agent's recent actions in the Agents tab to understand what's happening.

5. The gate is adding latency to my agent. Is this normal?

Gate evaluation takes approximately 10ms. This is normal and by design — your agent is trading 10ms for the guarantee that bad actions are caught. If latency is higher than expected, check your network connection to GovernHQ or contact support.